

dl-lab-assignment-4

December 14, 2023

```
[ ]: ! pip install tensorflow
```

Q1. Implement RNN, LSTM and GRU for MNIST dataset.

1) *RNN* :-

```
[24]: import torch
import torchvision
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.utils.data import DataLoader
import torchvision.datasets as datasets
import torchvision.transforms as transforms
from tensorflow.keras.datasets import mnist
```

```
[26]: #Setting up CPU or GPU
device=torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
[27]: #Setting hyperparameters
input_size=28
sequence_length=28
num_layers=3
hidden_size=256
num_classes=10
learning_rate=0.001
batch_size=64
num_epochs=5
```

```
[28]: #Defining the model-init and forward
class RNN(nn.Module):
    def __init__(self,input_size,hidden_size,num_layers,num_classes):
        super(RNN,self).__init__()
        self.hidden_size=hidden_size
        self.num_layers=num_layers
        self.rnn=nn.RNN(input_size,hidden_size,num_layers,batch_first=True)
        self.fc=nn.Linear(hidden_size*sequence_length,num_classes)
```

```

def forward(self,x):
    h0=torch.zeros(self.num_layers,x.size(0),self.hidden_size).to(device)
    out, _ =self.rnn(x,h0)
    out=out.reshape(out.shape[0],-1)
    out=self.fc(out)
    return out

```

```

[29]: #Getting the dataset
train_dataset=datasets.MNIST(root='dataset/',train=True,transform=transforms.
    ↳ToTensor(),download=True)
test_dataset=datasets.MNIST(root='dataset/',train=False,transform=transforms.
    ↳ToTensor(),download=True)

```

```

[30]: train_loader=DataLoader(dataset=train_dataset,batch_size=batch_size,shuffle=True)
test_loader=DataLoader(dataset=test_dataset,batch_size=batch_size,shuffle=True)

```

```

[15]: #Call the model
model=RNN(input_size,hidden_size,num_layers,num_classes)

```

```

[16]: #Define Loss function and Optimization Algorithm
criterion=nn.CrossEntropyLoss()
optimizer=optim.Adam(model.parameters(),lr=learning_rate)

```

```

[17]: #Training the data
for epoch in range(num_epochs):
    for batch_idx, (data,targets) in enumerate(train_loader):
        data=data.to(device=device).squeeze(1)
        targets=targets.to(device=device)

        scores=model(data)
        loss=criterion(scores,targets)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

```

```

[19]: model.eval()

```

```

[19]: RNN(
    (rnn): RNN(28, 256, num_layers=3, batch_first=True)
    (fc): Linear(in_features=7168, out_features=10, bias=True)
)

```

```

[20]: def check_aacuracy(loader,model):
    if loader.dataset.train:
        print("Accuracy of Training data")
    else:
        print("Accuracy of Testing data")

```

```

num_correct=0
num_samples=0
model.eval()
with torch.no_grad():
    for x,y in loader:
        x=x.to(device=device).squeeze(1)
        y=y.to(device=device)

        scores=model(x)
        _,predictions=scores.max(1)
        num_correct=num_correct+(predictions==y).sum()
        num_samples=num_samples+predictions.size(0)
print((float(num_correct)/float(num_samples))*100)
model.train()

```

```

[21]: check_aaccuracy(train_loader,model)
      check_aaccuracy(test_loader,model)

```

Accuracy of Training data
97.12833333333334
Accuracy of Testing data
96.66

2)LSTM:-

```

[31]: import torch
      import torchvision
      import torch.nn as nn
      import torch.optim as optim
      import torch.nn.functional as F
      from torch.utils.data import DataLoader
      import torchvision.datasets as datasets
      import torchvision.transforms as transforms

```

```

[32]: #Setting up CPU or GPU
      device=torch.device('cuda' if torch.cuda.is_available() else 'cpu')

```

```

[33]: #Setting hyperparametrs
      input_size=28
      sequence_length=28
      num_layers=2
      hidden_size=256
      num_classes=10
      learning_rate=0.001
      batch_size=64
      num_epochs=2

```

```
[34]: #Defining the model-init and forward
class LSTM(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers, num_classes):
        super(LSTM, self).__init__()
        self.hidden_size=hidden_size
        self.num_layers=num_layers
        self.lstm=nn.LSTM(input_size, hidden_size, num_layers, batch_first=True)
        self.fc=nn.Linear(hidden_size*sequence_length, num_classes)

    def forward(self, x):
        h0=torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(device)
        c0=torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(device)
        out, _ =self.lstm(x, (h0, c0))
        out=out.reshape(out.shape[0], -1)
        out=self.fc(out)
        return out

[35]: #Getting the dataset
train_dataset=datasets.MNIST(root='dataset/', train=True, transform=transforms.
    ↳ToTensor(), download=True)
test_dataset=datasets.MNIST(root='dataset/', train=False, transform=transforms.
    ↳ToTensor(), download=True)

[36]: train_loader=DataLoader(dataset=train_dataset, batch_size=batch_size, shuffle=True)
test_loader=DataLoader(dataset=test_dataset, batch_size=batch_size, shuffle=True)

[37]: #Call the model
model=LSTM(input_size, hidden_size, num_layers, num_classes)

[38]: #Define Loss function and Optimization Algorithm
criterion=nn.CrossEntropyLoss()
optimizer=optim.Adam(model.parameters(), lr=learning_rate)

[39]: #Training the data
for epoch in range(num_epochs):
    for batch_idx, (data, targets) in enumerate(train_loader):
        data=data.to(device=device).squeeze(1)
        targets=targets.to(device=device)

        scores=model(data)
        loss=criterion(scores, targets)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

[40]: model.eval()
```

```
[40]: LSTM(
      (lstm): LSTM(28, 256, num_layers=2, batch_first=True)
      (fc): Linear(in_features=7168, out_features=10, bias=True)
    )
```

```
[41]: def check_aacuracy(loader,model):
      if loader.dataset.train:
          print("Accuracy of Training data")
      else:
          print("Accuracy of Testing data")
      num_correct=0
      num_samples=0
      model.eval()
      with torch.no_grad():
          for x,y in loader:
              x=x.to(device=device).squeeze(1)
              y=y.to(device=device)

              scores=model(x)
              _,predictions=scores.max(1)
              num_correct=num_correct+(predictions==y).sum()
              num_samples=num_samples+predictions.size(0)
          print((float(num_correct)/float(num_samples))*100)
      model.train()
```

```
[42]: check_aacuracy(train_loader,model)
      check_aacuracy(test_loader,model)
```

```
Accuracy of Training data
98.20166666666667
Accuracy of Testing data
97.92999999999999
```

3)GRU:-

```
[43]: import torch
      import torchvision
      import torch.nn as nn
      import torch.optim as optim
      import torch.nn.functional as F
      from torch.utils.data import DataLoader
      import torchvision.datasets as datasets
      import torchvision.transforms as transforms
```

```
[44]: #Setting up CPU or GPU
      device=torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
[45]: #Setting hyperparameters
input_size=28
sequence_length=28
num_layers=2
hidden_size=256
num_classes=10
learning_rate=0.001
batch_size=64
num_epochs=2
```

```
[46]: #Defining the model-init and forward
class GRU(nn.Module):
    def __init__(self,input_size,hidden_size,num_layers,num_classes):
        super(GRU,self).__init__()
        self.hidden_size=hidden_size
        self.num_layers=num_layers
        self.gru=nn.GRU(input_size,hidden_size,num_layers,batch_first=True)
        self.fc=nn.Linear(hidden_size*sequence_length,num_classes)

    def forward(self,x):
        h0=torch.zeros(self.num_layers,x.size(0),self.hidden_size).to(device)
        out, _ =self.gru(x,h0)
        out=out.reshape(out.shape[0],-1)
        out=self.fc(out)
        return out
```

```
[47]: #Getting the dataset
train_dataset=datasets.MNIST(root='dataset/',train=True,transform=transforms.
    ↳ToTensor(),download=True)
test_dataset=datasets.MNIST(root='dataset/',train=False,transform=transforms.
    ↳ToTensor(),download=True)
```

```
[48]: train_loader=DataLoader(dataset=train_dataset,batch_size=batch_size,shuffle=True)
test_loader=DataLoader(dataset=test_dataset,batch_size=batch_size,shuffle=True)
```

```
[49]: #Call the model
model=GRU(input_size,hidden_size,num_layers,num_classes)
```

```
[50]: #Define Loss function and Optimization Algorithm
criterion=nn.CrossEntropyLoss()
optimizer=optim.Adam(model.parameters(),lr=learning_rate)
```

```
[51]: #Training the data
for epoch in range(num_epochs):
    for batch_idx, (data,targets) in enumerate(train_loader):
        data=data.to(device=device).squeeze(1)
        targets=targets.to(device=device)
```

```

scores=model(data)
loss=criterion(scores,targets)
optimizer.zero_grad()
loss.backward()
optimizer.step()

```

```
[52]: model.eval()
```

```

[52]: GRU(
      (gru): GRU(28, 256, num_layers=2, batch_first=True)
      (fc): Linear(in_features=7168, out_features=10, bias=True)
    )

```

```

[53]: def check_aacuracy(loader,model):
      if loader.dataset.train:
          print("Accuracy of Training data")
      else:
          print("Accuracy of Testing data")
      num_correct=0
      num_samples=0
      model.eval()
      with torch.no_grad():
          for x,y in loader:
              x=x.to(device=device).squeeze(1)
              y=y.to(device=device)

              scores=model(x)
              _,predictions=scores.max(1)
              num_correct=num_correct+(predictions==y).sum()
              num_samples=num_samples+predictions.size(0)
          print((float(num_correct)/float(num_samples))*100)
      model.train()

```

```

[54]: check_aacuracy(train_loader,model)
      check_aacuracy(test_loader,model)

```

```

Accuracy of Training data
98.45166666666667
Accuracy of Testing data
98.1

```